ITRI 625

Pfleeger Chapter 6

Whitman & Mattord Chapters 6 & 7

Environment of use

- The biggest difference between a network and a stand-alone device is the environment in which each operates.

- Networks can be described by several typical characteristics:

  - Anonymity. A network removes most of the clues by which we recognize acquaintances.
  - Automation. In some networks most points involved might be machines with only minimal human supervision.
  - Distance. Networks connect endpoints with such speed that users typically cant tell whether a remote site is near or far.
  - Opaqueness. Users cannot distinguish where the node they are connected to is located. They cant even tell if they are communicating with the same host than the previous time.
  - Routing diversity. The same interaction might follow different paths each time it is invoked.

What makes a network vulnerable?

- Anonymity. An attack can be mounted without ever coming into contact with the system, its administrators, or users.
- Many points of attack – both targets and origins. An attack can come from any host to any host, therefore many points of vulnerability.
- Sharing. Access is afforded to more users and systems. Access controls for single systems may be inadequate in networks.
- Complexity of system. A network OS is more complex than an OS for a single computing system. This inherent complexity makes it extremely difficult, if not impossible, to secure.
- Unknown perimeter. A network's expandability also implies uncertainty about the network boundary.
- Unknown path. There are many paths from one host to another. Network users seldom have control over the routing of their messages.

Main types of attacks

- Reconnaissance. For example port scans, for intelligence, for social engineering, etc.
- Threats in transit: Eavesdropping and Wiretapping. These types of threats can occur across all transmission media.
- Impersonation. For example authentication foiled by guessing, nonexistent authentication, masquerade, etc.
- Message confidentiality threats. For example misdelivery (although uncommon), exposure, traffic flow analysis, etc.
- Message integrity threats. For example falsification of messages and noise.

Main types of attacks *cont.*

- Format failures. For example malformed packets, protocol failures and implementation flaws.
- Web site vulnerabilities. For example web site defacement, application code errors, etc.
- Denial of service. For example transmission failure, connection flooding, traffic redirection, etc.
- Threats in active or mobile code. For example cookies, scripts, etc.
- Complex attacks. For example script attacks and building blocks.

Network security controls

- Architecture. Planning can be the strongest control, especially when planning to build in security as one of the key constructs.
    - Segmentation. Segmentation reduces the number of threats, and it limits the amount of damage a single vulnerability can allow.
    - Redundancy. Allowing a function to be performed on more than one node, to avoid "putting all your eggs in one basket".
    - Single points of failure. Is there a single point in the network that if it were to fail, could deny access to all or a significant part of the network?
    - Mobile agents. Mobile code and hostile agents are potential methods of attack, but they can also be forces for good.

Network security controls *cont.*

- Encryption. Encryption is probably the most important and versatile tool for a network security expert.
  - Link encryption. Data are encrypted just before the system places them on the physical communications link.
  - End-to-end encryption. Provides security from one end of a transmission to the other.
  - SSH encryption. Provides an authenticated and encrypted path to the shell or operating system command interpreter.
  - SSL encryption. SSL interfaces between applications and the TCP/IP protocols to provide server authentication, optional client authentication, and an encrypted communications channel between client and server.

Network security controls *cont.*

- Strong authentication. Authentication may be more difficult to achieve securely because of the possibility of eavesdropping and wiretapping.
  - One-time password. Good for one use only.
  - Challenge-response systems. A more sophisticated one-time password scheme. Uses a device that functions as an intermediary for authentication.
  - Digital distributed authentication. Developed due to the need to authenticate nonhuman entities in a computing system.
  - Kerberos. Used for authentication between intelligent processes, such as client-to-server tasks, or a user's workstation to other hosts (see p. 461). Supports authentication in distributed systems.

Network security controls *cont.*

- Wireless security. Being so exposed, it requires special measures to protect communication.
  - SSID. The service set identifier is the identification of an access point.
  - WEP. Meant to provide users privacy equivalent to that of a dedicated wire.
  - WPA and WPA2. Addresses the security deficiencies known in WEP.
  - Alarms and alerts. Devices that are placed inside protected networks to monitor what occurs inside the networks.
  - Honeypots. Computer systems open to attackers with the ideal of catching them.

Firewalls

- A firewall is a device that filters all traffic between a protected or "inside" network and a less trustworthy or "outside" network. Its purpose is to keep "bad" things outside a protected environment.

- Types of firewalls include:

1. Packet filtering gateways. The simplest type of firewall that controls access to packets on the basis of packet address or specific transport protocol type.

2. Stateful inspection firewalls. An improvement through maintaining state information from one packet to another in the input stream.

Firewalls *cont.*

3. Application proxies. A firewall that simulates the proper effects of an application so that the application receives only requests to act properly.

4. Circuit gateway. A firewall that essentially allows one network to be an extension of another

5. Guards. A sophisticated firewall that decides what services to perform on the user's behalf in accordance with its available knowledge, such as whatever it can reliably know of the outside user's identity, previous interactions, and so forth.

6. Personal firewalls. An application program that runs on a workstation to block unwanted traffic, usually from the network. Complements the work of a conventional firewall.

See table 6.6 on page 468 for a summary.

Intrusion detection systems

- A device, typically another separate computer, that monitors activity to identify malicious or suspicious events.

- Types of IDSs:

  - Signature based. Performs simple pattern matching and reports situations that match a pattern corresponding to a known attack type.

  - Heuristic based. Build a model of acceptable behavior and flag exceptions to that model.

  - Network based. A standalone device attached to the network to monitor traffic throughout that network.

  - Host based. Runs on a single workstation or client or host, to protect that one host.

Goals for IDSs

- An IDS should be fast, simple, and accurate, while at the same time being complete.

- An IDS could use some or all of the following design approaches:
  - Filter on packet headers.
  - Filter on packet content.
  - Maintain connection state.
  - Use complex, multipacket signatures.
  - Use minimal number of signatures with maximum effect.
  - Filter in real time, online.
  - Hide its presence.
  - Use optimal sliding time window size to match signatures.

Security for e-mail

- Threats to e-mails include:
  - Message interception (confidentiality).
  - Message interception (blocked delivery).
  - Message interception and subsequent replay.
  - Message content modification.
  - Message origin modification.
  - Message content forgery by outsider.
  - Message origin forgery by outsider.
  - Message content forgery by recipient.
  - Message origin forgery by recipient.
  - Denial of message transmission.

Security for e-mail *cont.*

- The requirements for secure e-mail include:

1. Message confidentiality.
2. Message integrity.
3. Sender authenticity.
4. Nonrepudiation.

- Encryption can be used to address all the requirements identified above. Not only the message is encrypted, but all the header information as well. A message integrity check can also be incorporated to further ensure integrity.

- Key management is however the major problem with e-mail encryption.

Questions

- Any questions?

ITRI 625

Pfleeger Chapter 7

Database security requirements

- **Physical database integrity**. The data in a database should be immune to physical problems, such as power failures. Someone should be able to reconstruct it if it is destroyed in a catastrophe.

- **Logical database integrity**. The structure of the database should be preserved. For example a modification in one field should not affect the other fields.

- **Element integrity**. The data contained in each element should be accurate. Can be ensured in three ways:

1. Field checks, where data type and format is checked.
2. Access control, where only certain people can make modifications.
3. Change log, where a log is kept of all changes made, thereby simplifying corrections.

Security requirements *cont.*

- **Auditability**. A record that is kept of all accesses that were made to the database, thereby assisting with the database's integrity.

- **Access control**. The database administrator specifies who should be allowed access to which data, at the view, relation, field, record or even element level.

- **User authentication**. The DBMS usually performs its own authentication, which in turn improves security and integrity.

- **Availability.** The DBMS tends to be taken for granted, and quite often "unavailability" is merely the system servicing another user.

Update integrity

- A system failure in the middle of modifying data is a serious problem (see page 514).

- The two-phase update has been developed to address this problem.

1. The intent phase. The DBMS gathers the resources it needs to complete the update, it prepares everything, but makes no changes. It is a repeatable phase and therefore no harm is done if a system failure occurs during this phase. The last event is committing, during which a commit flag is written. After committing the DBMS starts making permanent changes.

2. Also a repeatable phase during which the changes are made. If a failure occurs, the system can repair it by repeating the phase. Once this phase is complete, the database is complete.

Inference

- A way to infer or derive sensitive data from non-sensitive data.
- Direct attacks. Where users try to determine values of sensitive fields by seeking them directly with queries that yield few records.
- Indirect attacks (see page 521). Seeks to infer a final result based on one or more intermediate statistical results. Only neutral data is released, identifying characteristics are removed.
- Tracker attacks (see page 524). A DBMS may conceal data when a small number of entries make up a large proportion of the data revealed. These attacks fool the manager into locating the desired data by using additional queries that produce small results.

Inference *cont.*

- Linear system vulnerability (see page 525). With a little logic algebra, and luck in the distribution of the database contents, it may be possible to construct a series of queries that returns results relating to several different sets.

- See pages 529 – 535 for controls for inference attacks.

- There are no perfect solutions to the inference problem. The three main approaches to controlling it are:

1. Suppress obviously sensitive information.
2. Track what the user knows.
3. Disguise the data.

Multilevel databases

- Data cannot always only be classified as sensitive or non-sensitive.
- Three characteristics of database security that emerge are:

1. The security of a single element may be different from the security of other elements of the same record or from other values of the same attribute.
2. Two levels are inadequate to represent some security situations.
3. The security of an aggregate may differ from the security of the individual elements.

Proposals for multilevel security

- **Partitioning.** Where the database is divided into separate databases each at its own level of sensitivity. It does however destroy the advantage of elimination of redundancy and improved accuracy.
- **Encryption**. Each level of sensitive data can be stored in a table encrypted under a key unique to the level of sensitivity.
- **Integrity lock**. Provides both integrity and limited access. Each data item consists of three pieces. The actual data item itself, a sensitivity label and a checksum. The checksum is computed to prevent unauthorized modification.
- **Sensitivity lock.** A combination of a unique identifier and the sensitivity level. Because the identifier is unique, each lock relates to one particular record.

Security in data mining

- In a largely automated way, data mining applications sort and search through data.
- They present probable relationships, but these are not necessarily cause-and-effect relationships.
- **Privacy and sensitivity**. Although summarized results are used, individual privacy can still suffer. It suffers from the same kinds of aggregation and inference found in databases.
- **Data correctness and integrity.** Connecting the dots refers to drawing conclusions from relationships between discrete bits of data. The correct data should be collected correctly.

Security in data mining *cont.*

- **Correcting mistakes in data.** Quite often correcting a mistake requires a mistake to be rectified on a master record. In typical marketing scenarios this tends not to be the case. Data mining often takes place across databases that do not have shared keys, so they use data fields as keys. This typically leads to a mistake occurring in several places.

- **Using comparable data.** Data semantics is very important in data mining. When comparing fields in two different databases, they should be in the same format, to avoid badly distorted statistics.

- **Eliminating false matches.** Data mining will inevitably produce false positives and missed connections. We need to be sensitive to the inherent inaccuracy of data mining approaches. Correctness of results and interpretations are major security issues.

Questions

- Any questions?

ITRI 625

Pfleeger Chapter 9

Privacy concepts

- Privacy is the right to control who knows certain aspects about you, your communications, and your activities.

- Information privacy has three aspects, sensitive data, affected parties, and controlled disclosure.

1. Controlled disclosure. As soon as you give out your number for example, your control is diminished because it depends in part on what someone else does. You do not control the other person.

2. Sensitive data. Some examples of data many people consider private include, identity, finances, legal matters, religion, etc.

3. Affected subject. The entity involved governs the data and possible consequences.

Computer related privacy problems

- The eight dimensions of computer related privacy are:

1. Information collection. Only with knowledge and explicit consent.
2. Information usage. Only for certain specified purposes.
3. Information retention. Only for a set period of time.
4. Information disclosure. Only to authorized people.
5. Information security. Appropriate mechanisms are used to ensure protection.
6. Access control. All modes of access are controlled.
7. Monitoring. Logs are maintained showing all accesses.
8. Policy changes. Less restrictive policies are never applied after-the-fact to already obtained data.

Steps to protect against privacy loss

- Several steps that any entity can take to help safeguard private data include:

- Data minimization.
- Data anonymization.
- Audit trail.
- Security and controlled access.
- Training.
- Quality.
- Restricted usage.
- Data left in place.
- Policy.

Authentication and privacy

- Authentication can refer to authenticating an individual, identity, or attribute.
- An individual is an unique person. There are relatively few ways of identifying an individual, and usually weak authentication is acceptable.
- An identity is a character string or similar descriptor. From a privacy standpoint, there may or may not be ways to connect different identities. Depending on the application, it can typically be done through linking, especially when authentication is not ideal and anonymity is required.
- An attribute is a characteristic. By linking various characteristics, privacy can quite often be jeopardized. Research has indicated that the process of effectively anonymizing data is extremely difficult.

Privacy on the web

- The internet is perhaps the greatest threat to privacy.
- It is like a nightmare of a big, unregulated bazaar, where every word you speak can be heard by many others.
- Payments on the web are perhaps one of the main privacy concerns, and nowadays mainly credit cards or payment schemes are used for online transactions.
- Sites and portals often require registering before being able to use their services. More often than not this information is used for marketing or to show that advertisements are warranted.
- Third party ads, contests and offers are great ways for companies to collect information and draw links.

Privacy on the web *cont.*

- Precautions for web surfing include limiting cookies and web bugs, two technologies that are frequently used to monitor a user's activities without the user's knowledge.
- Cookies are files of data sent by a website, and are a cheap way of transferring storage needs from a website to a user.
- Third party cookies are for organizations other than the webpage's owner.
- A cookie is a tracking device, whereas a web bug is an invisible image that invites or invokes a process. The image is typically 1 x 1 pixel, so virtually invisible on modern resolutions.

Privacy on the web *cont.*

- Spyware code is designed to spy on a user to obtain information.
- Keystroke loggers, the computer equivalent of a telephone wiretap.
- Hijackers, software that hijacks a program installed for a different purpose.
- Adware, displays selected ads in pop-up windows with the aim of obtaining personal information.
- Drive-by installation, a means of tricking a user into installing software. It conceals from the user the real code being installed.

Impacts on emerging technologies

- Applications of three emerging technologies have inherent risks for privacy.

1. Radio frequency identification (RFID). Small, low power wireless radio transmitters. When a tag receives a signal it sends its ID number in response. Its major concerns are the ability to track people wherever, and correctness.

2. Electronic voting. Ensuring anonymity whilst voting using computers is extremely difficult. Both in capturing the vote and in transmitting it to the election headquarters.

3. Voice over IP (VoIP). Even if solidly encrypted, the source and destination of the phone call will be somewhat exposed through packet headers.

4. Internet of Things (IoT). Insecure products can lead to potentially catastrophic consequences.

Questions

- Any Questions?